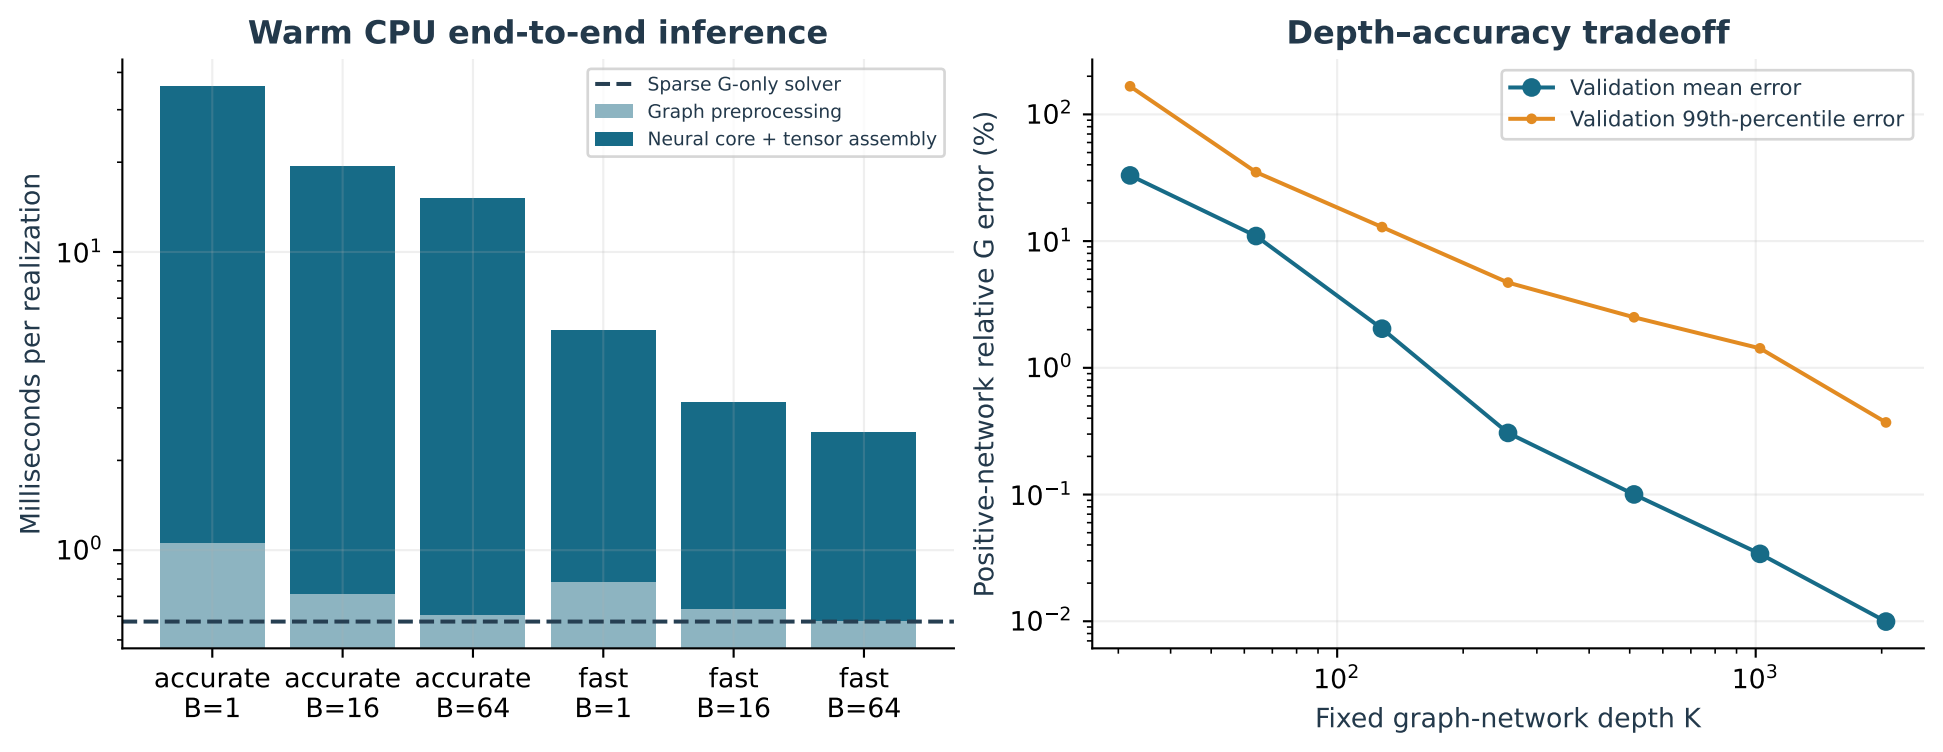
Warm CPU end-to-end inference
Sparse G-only solver
Graph preprocessing
Neural core + tensor assembly
Milliseconds per realization
10^1
10^0
accurate B=1
accurate B=16
accurate B=64
fast B=1
fast B=16
fast B=64
Depth–accuracy tradeoff
Validation mean error
Validation 99th-percentile error
Positive-network relative G error (%)
10^2
10^1
10^0
10^{-1}
10^{-2}
10^2
10^3
Fixed graph-network depth K